

Markov Decision Processes (MDPs)

Ron Parr
CSCI 2951-F
Brown University

With thanks to Kris Hauser for some slides

The MDP Framework

- State space: S
- Action space: A
- Transition function: P
- Reward function: $R(s,a,s')$ or $R(s,a)$ or $R(s)$
- Policy: $\pi(s) \rightarrow a$
- Discount factor: γ

Objective: *Maximize expected, discounted return
(sum of rewards)*

Applications of MDPs

- AI/Computer Science
 - Robotic control
(Koenig & Simmons, Thrun et al., Kaelbling et al.)
 - Air Campaign Planning (Meuleau et al.)
 - Elevator Control (Barto & Crites)
 - Computation Scheduling (Zilberstein et al.)
 - Control and Automation (Moore et al.)
 - Spoken dialogue management (Singh et al.)
 - Cellular channel allocation (Singh & Bertsekas)

Applications of MDPs

- Economics/Operations Research
 - Fleet maintenance (Howard, Rust)
 - Road maintenance (Golabi et al.)
 - Packet Retransmission (Feinberg et al.)
 - Nuclear plant management (Rothwell & Rust)
 - Debt collection strategies (Abe et al.)
 - Data center management (DeepMind)

Applications of MDPs

- EE/Control
 - Missile defense (Bertsekas et al.)
 - Inventory management (Van Roy et al.)
 - Football play selection (Patek & Bertsekas)
- Agriculture
 - Herd management (Kristensen, Toft)
- Other
 - Sports strategies
 - Board games
 - Video games

The Markov Assumption

- Let S_t be a random variable for the state at time t
- $P(S_t | A_{t-1}S_{t-1}, \dots, A_0S_0) = P(S_t | A_{t-1}S_{t-1})$
- Markov is special kind of *conditional independence*
- Future is independent of past given current state, **action** (similar to HMM assumptions, but adds actions)

About Rewards

- $R(s,a,s')$ is most general – typically interpreted to associate rewards with transitions. Reward is accrued in state s .
- $R(s,a)$ – Any $R(s,a,s')$ model can be converted to this w/o changing the optimal policy (because of linearity of expectation)
- $R(s)$ – Simplest to write and work with. In general, cannot convert from $R(s,a)$ w/o changing the optimal policy.
- Can always convert from less complicated reward models to more complicated (upwards in this list) w/o consequences

Understanding Discounting: $0 \leq \gamma \leq 1$

- Mathematical motivation
 - Keeps values bounded
 - What if I promise you \$0.01 every day you visit me?
- Economic motivation
 - Discount comes from inflation
 - Promise of \$1.00 in future is worth \$0.99 today
- Probability of dying (losing the game)
 - Suppose ϵ probability of dying at each decision interval
 - Transition w/prob ϵ to state with value 0
 - Equivalent to $1 - \epsilon$ discount factor

Discounting in Practice


- Often chosen unrealistically low
 - Faster convergence of the algorithms we'll see later
 - Leads to slightly myopic policies
- Can reformulate most algs. for avg. reward
 - Mathematically uglier
 - Somewhat slower run time

Value Determination

Determine the value of each state under policy π

$$V^\pi(s) = R(s, \pi(s)) + \gamma \sum_{s'} P(s' | s, \pi(s)) V^\pi(s')$$

Bellman Equation for a fixed policy π



$$V^\pi(s_1) = 1 + \gamma(0.4V^\pi(s_2) + 0.6V^\pi(s_3))$$

Matrix Form

$$\mathbf{P}^\pi = \begin{pmatrix} P(s_1 | s_1, \pi(s_1)) & P(s_2 | s_1, \pi(s_1)) & P(s_3 | s_1, \pi(s_1)) \\ P(s_1 | s_2, \pi(s_2)) & P(s_2 | s_2, \pi(s_2)) & P(s_3 | s_2, \pi(s_2)) \\ P(s_1 | s_3, \pi(s_3)) & P(s_2 | s_3, \pi(s_3)) & P(s_3 | s_3, \pi(s_3)) \end{pmatrix}$$

$$\mathbf{V}^\pi = \gamma \mathbf{P}^\pi \mathbf{V}^\pi + \mathbf{R}^\pi$$

Generalization of the game show example from earlier

How to solve this system efficiently? Does it even have a solution?

Solving for Values

$$\mathbf{V}^\pi = \gamma \mathbf{P}^\pi \mathbf{V}^\pi + \mathbf{R}^\pi$$

For moderate numbers of states we can solve this system exactly:

$$\mathbf{V}^\pi = \underbrace{(\mathbf{I} - \gamma \mathbf{P}^\pi)^{-1}} \mathbf{R}^\pi$$

Guaranteed invertible because $\gamma \mathbf{P}^\pi$
has spectral radius < 1 when $\gamma < 1$

Iteratively Solving for Values

$$\mathbf{V}^\pi = \gamma \mathbf{P}^\pi \mathbf{V}^\pi + \mathbf{R}^\pi$$

For larger numbers of states we can solve this system *indirectly*:

$$\mathbf{V}^\pi_{i+1} = \gamma \mathbf{P}^\pi \mathbf{V}^\pi_i + \mathbf{R}^\pi$$

Guaranteed convergent because $\gamma \mathbf{P}^\pi$
has spectral radius < 1

Converges to V^π , which we call a fixed point because updates
Don't change the value any more

When to stop an iterative solver?

- Just pick some big number of iterations
- Use convergence rates to bound minimum number of iterations required to get within some range of true value function
- Dynamically decide when to stop when change in value function is small from one iteration to the next (also can be guided by theory)

Interpreting the Iterations

- Suppose $V^{\pi}_0 = 0$, and R is defined on (s,a)
- Then $V^{\pi}_1 = R^{\pi}$ (value of executing 1 step of π)
- $V^{\pi}_2 = R^{\pi} + \gamma P^{\pi} V^{\pi}_1 = R^{\pi} + \gamma P^{\pi} R^{\pi}$
(expected value of executing 2 steps of π)
- $V^{\pi}_3 = R^{\pi} + \gamma P^{\pi} V^{\pi}_2 = R^{\pi} + \gamma P^{\pi} R^{\pi} + \gamma^2 (P^{\pi})^2 R^{\pi}$
(expected value of executing 3 steps of π)
- Can interpret these as the value of a **finite horizon** problem, where everything **stops** after i steps

Interpretation Continued

- $V^{\pi}_{\infty} = (I - \gamma P^{\pi})^{-1} R = V^{\pi}$ = infinite horizon values
- Infinite horizon = value of running π forever
- Nota bene: This **interpretation** applies when $V^{\pi}_0 = 0$, but iteration converges to V^{π} for any choice of V^{π}_0

Notation Alert

- Policy (π) is sometimes used as a subscript rather than superscript by some authors
- π may be dropped if there (should be) no confusion about which policy is under evaluation
- Some authors (e.g., textbook) use $T(s,a,s')$ for $P(s' | s,a)$
- Most CS authors use V for the value function
 - textbook uses U
 - Some from operations research use J or other letters

Establishing Convergence

- Eigenvalue analysis
- Monotonicity
 - Assume all values start pessimistic
 - One value must always increase
 - Can never overestimate
 - Easy to prove
- Contraction analysis...
(Proof included but not discussed in interest of time)

Contraction Analysis

- Define maximum norm

$$\|V\|_{\infty} = \max_i |V[i]|$$

- Consider two value functions V^a and V^b each at iteration 1:

$$\|V_1^a - V_1^b\|_{\infty} = \varepsilon$$

- WLOG say

$$V_1^a \leq V_1^b + \vec{\varepsilon} \quad (\text{Vector of all } \varepsilon\text{'s})$$

Contraction Analysis Contd.

- At next iteration for V^b :

$$V_2^b = R + \gamma P V_1^b$$

- For V^a

$$V_2^a = R + \gamma P(V_1^a) \leq R + \gamma P(V_1^b + \vec{\varepsilon}) = R + \gamma P V_1^b + \gamma P \vec{\varepsilon} = R + \gamma P V_1^b + \gamma \vec{\varepsilon}$$

- Conclude:

Distribute

$$\|V_2^a - V_2^b\|_{\infty} \leq \gamma \varepsilon$$


Importance of Contraction

- Any two value functions get closer
- True value function V^* is a fixed point (value doesn't change with iteration)
- Max norm distance from V^* decreases *dramatically* quickly with iterations

$$\|V_0 - V^*\|_{\infty} = \varepsilon \rightarrow \|V_n - V^*\|_{\infty} \leq \gamma^n \varepsilon$$

Finding Good Policies

Suppose an expert told you the “true value” of each state:



Improving Policies

- How do we get the optimal policy?
- If we knew the values under the optimal policy, then just take the optimal action in every state
- How do we define these values?
- **Fixed point** equation with choices (Bellman equation):

$$V^*(s) = \max_a R(s,a) + \gamma \sum_{s'} P(s'|s,a) V^*(s')$$

Decision theoretic optimal choice given V^*
If we know V^* , picking the optimal action is easy
If we know the optimal actions, computing V^* is easy
How do we compute both at the same time?

Value Iteration


We can't solve the system directly with a max in the equation
Can we solve it by iteration?


$$V_{i+1}(s) = \max_a R(s,a) + \gamma \sum_{s'} P(s'|s,a) V_i(s')$$

- Called *value iteration* or simply *successive approximation*
- Same as value determination, but we can *change* actions
- Converges to V^*
- Convergence:
 - Can't do eigenvalue analysis (not linear)
 - Still monotonic
 - Still a contraction in max norm (fun exercise)
 - Converges quickly


Robot Navigation Example

(from Russell & Norvig AIMA text)




- The robot (shown ) lives in a world described by a 4x3 grid of squares with square (2,2) occupied by an obstacle
- A state is defined by the square in which the robot is located: (1,1) in the above figure
→ 11 states

From Burgard et al.,
"Experiences with an interactive museum tour-guide robot"



Action (Transition) Model




U brings the robot to:

- (1,2) with probability 0.8
- (2,1) with probability 0.1
- (1,1) with probability 0.1

- In each state, the robot's possible actions are {U, D, R, L}
 - For each action:
 - With probability 0.8 the robot does the right thing (moves up, down, right, or left by one square)
 - With probability 0.1 it moves in a direction perpendicular to the intended one
 - If the robot can't move, it stays in the same square
- [This model satisfies the Markov condition]

Action (Transition) Model




L brings the robot to:

- (1,1) with probability $0.8 + 0.1 = 0.9$
- (1,2) with probability 0.1

- In each state, the robot's possible actions are {U, D, R, L}
 - For each action:
 - With probability 0.8 the robot does the right thing (moves up, down, right, or left by one square)
 - With probability 0.1 it moves in a direction perpendicular to the intended one
 - If the robot can't move, it stays in the same square
- [This model satisfies the Markov condition]

Terminal States, Rewards, and Costs




"terminal" states
Not part of formal
MDP specification.
Usually handled by
forcing state to have a
fixed value, e.g. +1

- Two terminal states: (4,2) and (4,3)
- Rewards:
 - $R(4,3) = +1$ [The robot finds gold]
 - $R(4,2) = -1$ [The robot gets trapped in quicksand]
 - $R(s) = -0.04$ in all other states
- This example (from the Russell & Norvig text) assumes no discounting ($\gamma=1$)
- Discussion: Is this a good modeling decision?

How to Implement Terminal States

- Modify your algorithm
 - For states s that are “terminal”
 - For an iterative solver, just set $V(s)=R(s)$ at each iteration
 - If using matrix inversion, hack your matrix
- Modify your MDP
 - Create a state T with $R(T)=0$, $P(T|T,a)=1$ for all a
 - For all states s that are “terminal”
 - Set $P(T|s,a) = 1$ for all a
 - This forces $V(s)=R(s)$


The Optimal Policy is *Stationary*



- A **stationary policy** is a complete map π : state \rightarrow action
- For each non-terminal state it recommends an action, independent of when and how the state is reached
- Under the Markov and infinite horizon assumptions, the optimal policy π^* is necessarily a stationary policy
[The best action in a state does not depend on the past]

Is it obvious which policy is optimal for this problem?


(Stationary) Policy




- A **stationary policy** is a complete map $\pi: \text{state} \rightarrow \text{action}$
- For each non-terminal state it recommends an action, independent of when and how the state is reached
- Under the Markov Decision Process, the optimal policy π^* is necessarily a stationary policy
[The best action in a state does not depend on the past]

The optimal policy tries to avoid
"dangerous" state (3,2)


Optimal Policies for Various $R(s)$




$R(s) = -0.04$



$R(s) = -2$



$R(s) = -0.01$



$R(s) > 0$

Bellman Equation

3	→	→	→	+1
2	↑		↑	-1
1	↑	←	←	←
	1	2	3	4

The utility of s depends on the utility of other states s' (possibly, including s), and vice versa

$Appl(s)$ used if not all actions are defined in all states

- If s is terminal:

$$V(s) = R(s)$$

- If s is non-terminal:

$$V(s) = R(s) + \max_{a \in Appl(s)} \sum_{s' \in Succ(s,a)} P(s'|s,a) V(s')$$

[Bellman equation]

- $\pi^*(s) = \arg \max_{a \in Appl(s)} \sum_{s' \in Succ(s,a)} P(s'|s,a) V(s')$

Value Iteration Applied

3	0	0	0	+1
2	0		0	-1
1	0	0	0	0
	1	2	3	4

→

3	0.81	0.87	0.92	+1
2	0.76		0.66	-1
1	0.71	0.66	0.61	0.39
	1	2	3	4

1. Initialize the utility of each non-terminal states to $V_0(s) = 0$
2. For $t = 0, 1, 2, \dots$ do

$$V_{t+1}(s) = R(s) + \max_{a \in Appl(s)} \sum_{s' \in Succ(s,a)} P(s'|s,a) V_t(s')$$


for each non-terminal state s

State Utilities/Values

3	0.81	0.87	0.92	+1
2	0.76		0.66	-1
1	0.71	0.66	0.61	0.39
	1	2	3	4

- The utility of a state s is the maximal expected amount of reward that the robot will collect from s and future states by executing some action in each encountered state, until it reaches a terminal state (**infinite horizon**)
- Under the Markov and infinite horizon assumptions, **the utility of s is independent of when and how s is reached**
[It only depends on the possible sequences of states after s , not on the possible sequences before s]

Convergence of Value Iteration



Properties of Value Iteration

- VI converges to V^* ($\| \cdot \|_\infty$ from V^* shrinks by γ factor each iteration)
- Converges to optimal policy
- Why? (Because we figure out V^* , optimal policy is argmax)
- Optimal policy is stationary
(i.e. Markovian – depends only on current state)

Policy Iteration

Greedy Policy Construction

Let's *name* the action that looks best WRT V :

$$\pi_v(s) = \arg \max_a R(s,a) + \gamma \underbrace{\sum_{s'} P(s'|s,a)V(s')}_{\text{Expectation over next-state values}}$$

$$\pi_v = \text{greedy}(V)$$

Bootstrapping: Policy Iteration


Idea: Greedy selection is useful even with suboptimal V

Guess $\pi_v = \pi_0$

V_{π_v} = value of acting on π_v
(solve linear system)

$\pi_v \leftarrow \text{greedy}(V_{\pi_v})$

Repeat until
policy doesn't
change



Guaranteed to find optimal policy

Usually takes very small number of iterations

Computing the value functions is the expensive part

Comparing VI and PI

- VI
 - Value changes at every step
 - Policy *may* change before exact value of policy is computed
 - Many relatively cheap iterations
- PI
 - Alternates policy/value updates
 - Solves for value of each policy *exactly*
 - Fewer, slower iterations (need to invert matrix)
- Convergence
 - Both are contractions in max norm
 - PI is *shockingly* fast (small number of iterations) in practice

Computational Complexity

- VI and PI are both contraction mappings w/rate γ
(we didn't prove this for PI in slides)
- VI costs less per iteration
- For n states, a actions PI tends to take $O(n)$ iterations in practice
 - Recent(ish) results indicate $\sim O(n^2 a / (1 - \gamma))$ worst case
 - Interesting aside: Biggest insight into PI came ~ 50 years after the algorithm was introduced

MDP Limitations → Reinforcement Learning

- MDP operate at the level of *states*
 - States = atomic events
 - We usually have exponentially (or infinitely) many of these
- We assume P and R are known
- Machine learning to the rescue!
 - Infer P and R (implicitly or explicitly from data)
 - Generalize from small number of states/policies

Bonus Material

A Unified View of Value Iteration and Policy Iteration

Notation

- Update for for a fixed policy – definition of T^π operator (matrix-vector form):

$$T^\pi V \equiv R^\pi + \gamma P^\pi V$$

- Update with policy improvement – definition of the T operator:

$$TV(s) \equiv \max_a r(s, a) + \gamma \sum_{s'} P(s'|s, a) V(s')$$

Value Determination

- For 0 steps $V_0 = R^\pi$
- For i steps $V_i = T^\pi V_{i-1} = T^\pi T^\pi V_{i-2} = \dots = (T^\pi)^i R^\pi$
- Infinite horizon $\lim_{i \rightarrow \infty} V_i = (T^\pi)^\infty R^\pi = (1 - \gamma P^\pi)^{-1} R^\pi = V^\pi$

Value Iteration (includes MAX)

- For 0 steps $V_0 = R$ (If R depends on a, pick a with the highest immediate reward)
- For i steps $V_i = TV_{i-1} = T^i R$
- Infinite horizon $\lim_{i \rightarrow \infty} V_i = T^\infty R = TV^* = V^*$

Modified Policy Iteration

- Guess V_0 (usually just R), and π - or set $\pi = \text{greedy}(V_0)$
- $i=1$
- Repeat until convergence*
 - For $j=1$ to n
 - $V_i = T^\pi V_{i-1}$
 - $i = i+1$
 - $\pi = \text{greedy}(V_{i-1})$
- Special cases: $n=1$ (VI), $n \rightarrow \infty$ (PI)

n steps of iterative
policy evaluation

Linear Programming Review

- Minimize: $\mathbf{c}^T \mathbf{x}$
- Subject to: $\mathbf{Ax} \geq \mathbf{b}$
- Can be solved in weakly polynomial time
- Arguably most common and important optimization technique in history

Linear Programming

$$V(s) = \max_a R(s,a) + \gamma \sum_{s'} P(s'|s,a) V(s')$$

Issue: Turn the non-linear max into a collection of linear constraints

$$\forall s,a : V(s) \geq \underbrace{R(s,a) + \gamma \sum_{s'} P(s'|s,a) V(s')}$$

MINIMIZE: $\sum_s V(s)$

Optimal action has
tight constraints

Weakly polynomial; slower than PI in practice
(though can be modified to behave like PI)